

One Network, Many Kinds of Evidence

A Critical Analysis of Multiomics Integration, Edge Evidence, and Hypothesis Generation in “The Molecular Human”

A critical analysis of “A roadmap to the molecular human linking multiomics with population traits and diabetes subtypes”

Halama, Zaghlool, Thareja, et al. (2024), Nature Communications, 15, 7111

Sheryl S.

Research Portfolio Series · Critical Analysis #02

September 2026

Table of Contents


1. Central Claim and Evidentiary Overview.....	3
2. Theoretical Foundation.....	4
2.1 Partial Correlation and Gaussian Graphical Models (GGM).....	4
2.2 Mutual Best Hits (MBH) as a Borrowed Bioinformatics Concept.....	4
2.3 Why GWAS, EWAS, and TWAS Are Not Interchangeable Evidence.....	4
2.4 Multiple-Testing Correction Across Different Comparison Families.....	5
2.5 Missing Data and Unequal Platform Overlap.....	5
3. Data and Study Design.....	5
4. Integration Methodology — The Central Critique.....	6
4.1 One Network, Many Kinds of Edges.....	6
4.2 The MBH and GGM Pipelines Are Not Preprocessed the Same Way.....	9
4.3 MBH as an Analogy to Orthology — What It Can and Cannot Claim.....	9
4.4 What the Network Establishes, Versus What It Would Take to Confirm a Mechanism.....	9
5. Case Studies as Illustrative Validation: Caffeine and HbA1c.....	10
6. GWAS, EWAS, and TWAS — Not Equivalent Evidence in Practice.....	11
7. Diabetes-Subtype Networks and Generalizability from 391 Participants.....	11
8. Strengths and Limitations.....	12
8.1 Strengths.....	12
8.2 Limitations.....	12
9. Reproducibility.....	14
10. My Critical Perspective.....	15
11. What Would Change My Mind.....	15
12. Potential Research Extensions.....	16
Extension 1 — Covariate-adjusted MBH as an alternative edge-selection rule.....	16
Extension 2 — Basophil-specific sensitivity of the TWAS lipid signal.....	16
Extension 3 — Reconstructing a small part of the COMics network and inspecting edge provenance.....	16
13. Questions the Paper Leaves Open.....	17
14. Lessons Learned.....	17
15. Connection to My Development as a Bioinformatician.....	17
16. Conclusion.....	18
References.....	19

1. Central Claim and Evidentiary Overview

Halama et al. report deep molecular phenotyping of 391 participants in the Qatar Metabolomics Study of Diabetes (QMDiab), combining 18 analytical platforms across blood, urine, and saliva into what they call “The Molecular Human”: an integrated network of more than 34,000 statistically significant trait-trait links spanning 6,304 molecular nodes. The paper's central claim is that this scale of integration — within-platform partial correlations (GGM), between-platform mutual best hits (MBH), and genome-, epigenome-, and transcriptome-wide associations (GWAS, EWAS, TWAS) — can be combined into a single network that is useful both for characterizing platform complementarity and for generating biological hypotheses, including about type 2 diabetes (T2D) subtypes. The authors make this network publicly explorable through a web tool, COMics.

What interested me most from a bioinformatics perspective was not any single finding but the integration step itself: several statistically and biologically different kinds of evidence (partial correlation, reciprocal best-hit correlation, and multiple genetics-, methylation-, and transcript-expression-based association analyses — GWAS, EWAS, and TWAS, further split into eQTL, meQTL, and eQTM sub-analyses) are drawn into one graph and, in the COMics interface, are visually distinguished mainly by the omics platform of the connected nodes rather than by the type or strength of evidence that produced the edge. Before working through the Methods section carefully, I would have treated an edge in this network as a roughly interchangeable unit of “association.” The paper's own methods make clear why that would be a mistake, and that distinction is what organizes the rest of this analysis.

Figure 1. Study and evidence-generation pipeline



Conceptual diagram constructed by the author; not reproduced from Halama et al. (2024).
Values and terminology drawn directly from the article text, abstract, and Figure 1 of the source paper.

Figure 1. Study and evidence-generation pipeline for Halama et al. (2024), from cohort and sample collection through to the integrated network and the COMics resource. Conceptual diagram constructed by the author; not reproduced from Halama et al. (2024).

2. Theoretical Foundation

This section covers only the concepts I needed to work through in order to follow my own critique below — it is not meant as a general multiomics primer.

2.1 Partial Correlation and Gaussian Graphical Models (GGM)

A simple pairwise (Pearson or Spearman) correlation between two molecular traits can be inflated by a third variable that drives both of them. A Gaussian Graphical Model instead estimates the partial correlation between every pair of traits within a platform, conditioning on all other traits measured on that same platform. This can remove some indirect associations transmitted through variables included in the model, but it does not prove that a retained edge represents a direct biological interaction — mediation through an unmeasured variable, or a shared upstream driver not included in the residualization, is still possible. In this paper, GGM edges were computed only within a platform, after residualizing each trait on age, sex, BMI, diabetes state, and genotype/proteome principal components, and after Bonferroni correction scaled to the number of traits on that platform.

2.2 Mutual Best Hits (MBH) as a Borrowed Bioinformatics Concept

The MBH approach used here is explicitly adapted from bidirectional best hit (BBH) methods in comparative genomics, where two genes from different species are called orthologs if each is the other's most similar match in a sequence search. Halama et al. apply the same reciprocity logic to statistical correlation: two traits on different platforms are linked only if each is the other's strongest Spearman correlation partner, out of all traits on the other platform. This is a hypothesis-free way to find between-platform relationships without specifying in advance which trait pairs should be compared, and the analogy to orthology is a genuinely useful way to think about what MBH is doing — finding a plausible best match, not confirming identity.

2.3 Why GWAS, EWAS, and TWAS Are Not Interchangeable Evidence

These analyses test statistical associations between genetic, methylation, or expression measurements and molecular traits, but the biological interpretation differs across layers. Germline genotype (GWAS) is fixed before the molecular phenotype is ever measured, so the molecular trait cannot cause the genotype. That removes one specific direction of reverse causation, though the association may still reflect linkage disequilibrium, population structure, horizontal pleiotropy, or an indirect biological pathway rather than a direct causal link, the responsible gene, or the true causal variant. DNA methylation (EWAS) and gene expression (TWAS), by contrast, are measured in the same blood draw as the phenotype and can plausibly be a consequence of the phenotype, a cause of it, or a shared downstream effect of something else (including, in blood, differences in circulating immune-cell composition). All three are compatible with a statistically significant cross-sectional association, so treating a GWAS hit, an EWAS hit, and a TWAS hit as equally strong evidence for a biological relationship is not something the underlying statistics support.

One terminology difference matters here. Halama et al. use “moTWAS” for associations between directly measured RNA expression (from RNA-sequencing) and other omics traits. This differs from the common contemporary use of “TWAS” for analyses based on genetically predicted gene expression (imputed from

genotype using a reference expression panel). Throughout this analysis, I use moTWAS according to the definition adopted by Halama et al. — an association with directly measured, not genetically predicted, transcript abundance.

2.4 Multiple-Testing Correction Across Different Comparison Families

Because the platforms contribute very different numbers of evaluated traits—from 36 total plasma N-glycan traits to 1,201 traits on the LD platform (Table 2 of the paper reports the number of traits actually evaluated per platform, which for LD is smaller than the 1,331 lipids described in the Methods as originally measured)—a single significance threshold based on an identical number of comparisons would not represent each platform pair equally. The authors use platform-pair-specific Bonferroni thresholds for MBH analyses and platform-specific Bonferroni thresholds for GGMs. They use 5×10^{-8} for the genetic association analyses and also adopt the same numerical threshold for the methylation- and transcript-expression-based association analyses, where the Methods describe it as an ad hoc significance level. The integrated network therefore combines edges retained under several different correction frameworks — platform-pair-sized Bonferroni for MBH, platform-sized Bonferroni for GGM, and a shared 5×10^{-8} cutoff of differing methodological status for the genetic-, methylation-, and expression-anchored analyses — which should be considered when comparing edge counts and network density across analysis types.

2.5 Missing Data and Unequal Platform Overlap

Not every participant has data on every platform (Table 2 of the paper reports between 229 and 356 shared samples per platform pair, out of 391 total), and within a platform, some molecular traits have substantial missingness. The GGM pipeline removed traits and then samples exceeding 50% missing values and imputed the remainder with k-nearest-neighbors; the MBH pipeline, as described in the Methods, was run on the available paired samples without an explicit missing-value-handling step described for correlation input beyond the pairwise sample overlap itself. I would want to test the stability of the reported edge counts under a different overlap threshold, but that is not something I can do from the material available to me.

3. Data and Study Design

Reconstructed as a pipeline: recruit a diabetes case-control cohort at a single hospital dermatology department in Doha, Qatar → collect non-fasting blood, urine, and saliva under identical protocols for cases and controls → profile the samples on 18 technically distinct platforms → compute within-platform GGMs, between-platform MBHs, and genotype/methylation/transcript-anchored associations → integrate all edges, together with GWAS-catalog lookups, into one network → build diabetes-subtype-specific sub-networks around previously published protein and metabolite signatures → release the whole network through the COMICS web interface.

A few cohort details matter for the rest of this analysis. QMDiab enrolled 391 participants with at least one omics measurement (193 female, 198 male; mean age 46.5 years, s.d. 12.9; mean BMI 29.7 kg/m², s.d. 6.0; 195 with T2D and 196 without), drawn from a multiethnic population of predominantly Arab, Filipino, and Indian background. Enrollment was continuous and not stratified by diabetes state, age, sex, BMI, or ethnicity, and participants were non-fasting at sample collection. The paper reports up to 6,304 quantitative molecular traits, 1,221,345 autosomal genetic variants after quality-control filtering, methylation at 470,837 CpG sites, and 57,942 transcriptomic traits (57,773 RNA transcripts plus 169 microRNAs).

Cross-platform correlation (MBH) analyses were computed using Spearman correlation on unscaled raw omics data, with platform-pairwise Bonferroni significance cutoffs. Within-platform GGM analyses instead used inverse-normal-scaled, covariate-residualized data (age, sex, BMI, diabetes state, and genotype/protein principal components), with per-platform Bonferroni cutoffs. The study also performed several association analyses connecting genetic variants, DNA methylation, RNA expression, and other omics traits, including eQTL, meQTL, eQTM, moQTL, moEWAS, and moTWAS analyses. These analyses used a fixed significance threshold of 5×10^{-8} , although the interpretation of that threshold and the resulting associations differs across molecular layers. This asymmetry in preprocessing between the MBH and GGM pipelines is a methodological detail I return to in Section 4.


4. Integration Methodology — The Central Critique

Every later stage of the paper — the case-study vignettes, the diabetes-subtype networks, the COMics resource itself — presents the integrated network as a single object to explore. None of those later stages re-annotates individual edges with the type or strength of evidence behind them once the edge has been added to the network. That is the throughline of this section.

4.1 One Network, Many Kinds of Edges

The source paper's main Results text and Fig. 1C report several statistically and biologically different kinds of network edge: GGM = 6,183, MBH = 2,103, meQTL = 1,381, eQTM = 15,991, eQTL = 17, EWAS/moEWAS = 3,772, TWAS/moTWAS = 1,660, and GWAS/moQTL, which is reported as 768 in the main Results text but as 786 in the “mo GWAS” summary box of Fig. 5A. Because these two values differ within the source article itself, I retain the distinction below (Figure 3) rather than silently picking one. Together with the phenotype-association tests (against age, sex, BMI, and T2D state), these are heterogeneous network-edge types — not a fixed count I want to assert precisely, since the integrated resource also includes GWAS-catalog lookups layered on top. Figure 2 sets several of the major categories side by side as a simplified conceptual grouping, not an exhaustive inventory of every association type in the network. What struck me is that in the COMics interface itself (Fig. 5A of the source paper), the legend distinguishes edges only as “positive/negative” and by data-type icon of the connected node (DNA, methylation, RNA, protein, and so on) — not by which of GGM, MBH, GWAS, EWAS, or TWAS produced the edge, and not by any measure of effect size or replication status. A reader exploring a single molecule's neighborhood in COMics would need to already know the underlying methods paper to know that a partial-correlation edge and a TWAS edge carry different evidentiary weight.

Figure 2. One integrated network, several different kinds of edge evidence



Conceptual diagram constructed by the author; not reproduced from Halama et al. (2024).

Figure 2. One integrated network, several different kinds of edge evidence — a simplified conceptual grouping of the major categories, not a complete count of every association type in the network. Conceptual diagram constructed by the author; not reproduced from Halama et al. (2024).




Chart reconstructed from values explicitly reported in the main text and Fig. 1C of Halama et al. (2024) (GGM = 6,183; MBH = 2,103; meQTL = 1,381; eQTM = 15,991; eQTL = 17; GWAS/moQTL = 768; EWAS/moEWAS = 3,772; TWAS/moTWAS = 1,660). Note: the source article reports two different GWAS/moQTL values — 768 in the main Results text (used here) versus 786 “mo GWAS” links in the Fig. 5A summary box. No values were estimated or interpolated.

Figure 3. Reported counts of each association/edge type integrated into the multiomics network. The GWAS/moQTL bar uses 768, the value stated in the main Results text of Halama et al. (2024); Fig. 5A of the source paper instead reports 786 “mo GWAS” links for the same category. All other values are reconstructed from the main text and Fig. 1C exactly as reported, with no values estimated or interpolated.

Table 1 summarizes what each of these edge/evidence types can and cannot support on its own, drawing on the distinctions worked through in this section and in Section 6.

Evidence type	What it measures	What it supports	What it cannot establish alone	Important limitation
GGM (partial correlation)	Direct within-platform statistical association, conditioned on other platform traits and on age/sex/BMI/diabetes/P Cs	A conditional statistical relationship between two traits after accounting for the other measured traits included in the model	A direct biochemical interaction, causal direction, or absence of mediation through an unmeasured variable	Only within-platform; cannot link across omics layers on its own
MBH (mutual best hit)	Reciprocal best-match Spearman correlation between traits on two different platforms, on raw unadjusted data	A hypothesis-free candidate link between platforms; sometimes molecular identity, sometimes pathway coupling	That the two linked traits are the same molecule; that the link is free of age/sex/BMI/diabetes confounding	Not covariate-adjusted in this pipeline (Section 4.2); ~28% of known-overlapping protein targets missed
GWAS / moQTL	Association between a genetic variant and a molecular trait, genome-wide significance threshold	Removes one direction of reverse causation (genotype cannot be caused by phenotype)	Full causality; the causal variant or gene; the mechanism linking variant and trait	Statistical significance does not imply the effect is large or clinically important
EWAS / moEWAS	Association between DNA methylation at a CpG site and a molecular trait	A candidate epigenetic correlate of the trait	Direction of causality; methylation could be cause, consequence, or shared downstream effect	Cross-sectional; blood-cell-composition differences can confound methylation associations
TWAS / moTWAS	Association between measured (not genetically predicted) gene expression in blood and a molecular trait	A candidate transcriptional correlate of the trait	Direction of causality; whether the source cell type drives the association	Adjusted for CD4/CD8 T cells, NK cells, B cells, monocytes, granulocytes, and eosinophils, but not separately for basophils — the cell type most linked to the top TWAS genes (Section 6)
Phenotype / disease association (e.g., variance explained)	How much of the variance in age, sex, BMI, or T2D state a platform's traits predict (random forest)	Which platforms most accurately track a phenotype overall	Which individual molecule drives the prediction; mechanism	A high variance-explained value can still reflect many small, non-specific contributions

Table 1. Evidence-type interpretation across the major edge types integrated into the Halama et al. (2024) network.

4.2 The MBH and GGM Pipelines Are Not Preprocessed the Same Way

This is the most concrete methodological point I could verify directly from the Methods section. The GGM (“Within-trait partial correlations”) pipeline explicitly residualizes each omics trait on age, sex, BMI, diabetes state, and genotype/protein principal components before computing partial correlations. The MBH (“Cross-platform correlations”) pipeline, as described, computes Spearman correlations directly on unscaled raw omics data, with no covariate adjustment step described. So a between-platform MBH edge could, in principle, be driven partly by a shared association with age, sex, BMI, or diabetes state that the within-platform GGM edges have already had removed — the two edge types in the same network were not held to the same confounding-control standard. I want to be careful here: the paper does not claim the two pipelines use identical preprocessing, and there may be a methodological reason I haven’t worked through — the reciprocity requirement of MBH could be sensitive to residualization in a way that made the authors choose not to residualize before computing it. I could not determine whether this was a deliberate design choice or simply a difference I am reading too much into. What I can say is that it limits how directly a GGM edge and an MBH edge in the same figure should be compared.


4.3 MBH as an Analogy to Orthology — What It Can and Cannot Claim

The bidirectional-best-hit analogy to orthology is, I think, a genuinely elegant framing, and the paper is careful to motivate it as a hypothesis-free way to link platforms rather than as proof of anything. But an MBH edge is a statement about reciprocal statistical association, not a statement that two measurements represent the same underlying molecule, and the paper’s own case studies illustrate a spectrum rather than a single relationship type. Some MBH edges connect what really does look like the same measured entity approached by two technologies — the paper reports 72–93% concordance between platforms measuring overlapping molecule sets (PM↔HDF, OLINK↔SOMA, IgG↔IgA). Others connect clearly distinct molecules linked by a substrate-product relationship (paraxanthine in saliva to 1,7-dimethylurate in urine, both downstream of caffeine) or by a binding relationship (thyroxine to its carrier protein SERPINA7). The paper is explicit that even between two proteomics platforms measuring overlapping protein targets, about 28% of common protein targets were not detected by MBH at all — a concrete number that shows MBH is an imperfect identity detector even in the best-case scenario where identity is actually the right interpretation. I interpret this as evidence that an MBH edge should be read case-by-case, informed by what the two connected platforms measure, rather than assumed by default to mean either “same molecule” or “direct biological interaction.”

4.4 What the Network Establishes, Versus What It Would Take to Confirm a Mechanism

Pulling Sections 4.1–4.3 together: the paper’s evidence chain reliably gets from an observed statistical association to a network edge to a biologically plausible hypothesis. Getting from a plausible hypothesis to a demonstrated mechanism requires additional steps — establishing temporal ordering, independent replication, and some form of perturbation or functional validation — that are outside what this study’s design was set up to provide, and the paper does not claim to have completed them. Figure 4 lays out that contrast explicitly, because I found it easy, on a first read, to mentally round “network edge” up to “confirmed relationship,” and the distinction is worth keeping visible going into the case studies and genetics- and expression-based findings discussed next.

Figure 4. What the reported associations establish, and what they do not



Conceptual diagram constructed by the author; not reproduced from Halama et al. (2024). Does not imply the authors claimed to complete the bottom-row validation stages.

Figure 4. What the reported associations establish, and what they do not. Conceptual diagram constructed by the author; not reproduced from Halama et al. (2024). Does not suggest the authors claimed to complete the validation stages in the bottom row.

5. Case Studies as Illustrative Validation: Caffeine and HbA1c

Before criticizing the network further, it is worth crediting what the caffeine and HbA1c vignettes (the paper's Fig. 2 and Supplementary Fig. 2A) actually demonstrate. Combining MBH edges across saliva, urine, and plasma metabolomics with within-platform GGM edges reconstructs a caffeine-metabolism pathway (caffeine → paraxanthine/theobromine/theophylline → further downstream metabolites) that matches independently published human caffeine pharmacokinetics. The HbA1c vignette similarly recovers known T2D-associated molecules (betaine, mannose) around a well-established clinical marker. These examples function as biological sanity checks, or illustrative validation cases: they demonstrate that the integration pipeline can recover selected relationships consistent with established biology. I don't think the paper declares these as formal, prespecified positive controls in the way a wet-lab experiment might use that term — they read to me more as worked examples the authors chose because the expected answer was already known, which is still a useful thing to show, just a narrower claim than a designed positive-control experiment would support.

This does not invalidate the approach, but it limits the strength of the interpretation that can be drawn from it. These sanity checks support the biological plausibility of the approach and demonstrate recovery of a handful of selected, already-known relationships — they do not estimate the network's sensitivity, specificity, precision, or false-discovery rate across the other tens of thousands of edges, most of which do not have decades of prior biochemistry to check against. I interpret the caffeine and HbA1c vignettes as evidence that the pipeline is doing something biologically sensible in at least these cases, not as evidence about the false-discovery rate among the novel, unreplicated associations the paper also reports (such as the ST3GAL1–glycan or 5-methyluridine–immune-gene findings discussed next).

6. GWAS, EWAS, and TWAS — Not Equivalent Evidence in Practice

The paper reports several genetics- and expression-based findings that are worth reading against the distinction laid out in Section 2.3. The GWAS finding linking variants near ST3GAL1 (a sialyltransferase gene) to IgA glycan traits is, in my reading, the strongest single result in the paper from a causal-inference standpoint: the genetic variant cannot be a consequence of glycan levels, ST3GAL1 has an annotated enzymatic role directly relevant to sialic acid transfer, and the association replicates the kind of genotype-to-molecular-trait link that GWAS is specifically suited to support.

The miRNA-methylation-genotype finding (three miRNAs associated with a CpG site that is itself associated with a nearby SNP) is presented as a novel CpG-miRNA axis. Because it combines a genetic anchor (rs45552131) with an epigenetic association (EWAS-type), the genetic layer again provides some protection against reverse causation for the SNP-CpG link specifically, but the CpG-miRNA association itself remains a cross-sectional association between two molecular layers measured in the same blood draw, and the paper does not report a temporal or interventional test of that specific link.

The TWAS results are where I think the causal ambiguity is most important, and the paper itself supplies part of the reason why. The majority of TWAS associations (67%, 1,114 of 1,660) were with lipids and lipoproteins, clustering around five gene transcripts, and the paper reports that four of those five genes (GATA2, HDC, MS4A3, FCER1A) are predominantly expressed by basophils according to the Human Protein Atlas. On first inspection, this pattern made the transcript-lipid associations look vulnerable to general immune-cell-composition confounding: if basophil-rich samples simply differ in their lipid profiles for unrelated reasons, an association with a basophil-marker gene could reflect cell mixture rather than a direct transcriptional link. The Methods, however, show that the moTWAS model is specified as “lm(OMICS ~ transcriptomics + AGE + SEX + BMI + DIAB + genoPC1 + genoPC2 + genoPC3 + CD8 + CD4 + NK + Bcell + Mono + Gran + Eos)” — meaning the model already adjusts for CD4 and CD8 T cells, NK cells, B cells, monocytes, granulocytes, and eosinophils. It would be inaccurate to describe this analysis as unadjusted for immune-cell composition; several major leukocyte populations are explicitly controlled for. The remaining concern is therefore narrower: basophil abundance was not listed as a separate covariate, even though four of the five transcripts dominating the reported lipid associations were described as predominantly basophil-expressed. Residual confounding from that one cell type specifically remains possible — the other adjustments would not remove it — but this is not something I can establish one way or the other from the reported analysis alone.

The Ingenuity Pathway Analysis (IPA) interaction diagram connecting GATA2, HDC, MS4A3, and FCER1A (Fig. 3D of the source paper) is a database-annotation step, drawing on previously curated interactions, rather than a new experiment performed by the authors in this study. It supports biological plausibility for these four genes belonging to a related pathway; it is not itself evidence generated from the QMDiab data.

7. Diabetes-Subtype Networks and Generalizability from 391 Participants

The diabetes-subtype networks (MARD, MOD, SIRD, SIDD, following the Ahlqvist et al. clustering framework) are built around protein and metabolite signatures that the paper's Methods and Discussion indicate were originally defined in a separate, earlier study of diabetes subtypes in the Qatar Biobank (QBB) cohort, and then overlapped with the present QMDiab multiomics dataset (Supplementary Data 16). Because I do not have access to that supplementary data table, I could not verify from the material available to me exactly how many of each subgroup's original signature molecules were present and matched within QMDiab, nor whether

QMDiab participants themselves were reclassified into MARD/MOD/SIRD/SIDD subgroups internally or whether the published signature list was applied directly. The supplementary material available with the uploaded document did not allow me to verify this point.

What is directly stated is that the MARD sub-network comprises 108 molecules measured across 11 platforms — a modest number relative to the roughly 6,300 total molecular traits in the full network, and the paper does not report the sample size within the 391-participant cohort that falls into each specific subgroup. Given that QMDiab is a single-site, non-fasting, multiethnic (predominantly Arab, Filipino, and Indian) cohort of 391 people, I would treat the subtype-specific molecular patterns described (for example, the osteomodulin-centered cardiovascular-calcification signature in MARD) as hypothesis-generating leads specific to this cohort and this subgroup-definition pipeline, not as claims that generalize automatically to T2D subtypes in other ancestries or clinical settings. The authors themselves are careful to frame the diabetes-subtype networks as offering “further insight” requiring “further investigation,” which I read as consistent with this caution.

I also want to flag, because the authors flag it themselves, that the cohort's non-fasting status and its mix of healthy and T2D participants means some of the reported associations — including, by their own account, part of the TWAS lipid signature — could be driven by molecular changes that are simply known features of T2D (elevated carbohydrates, lipids, and branched-chain amino acids) rather than a novel discovery. This is a limitation the paper states plainly rather than one I am inferring.

8. Strengths and Limitations

8.1 Strengths

- 18 technically distinct platforms measured on a single, deeply characterized cohort of 391 individuals is, on the numbers reported, an unusually deep degree of molecular phenotyping (with platform-pair overlap of 229–356 participants), and it allows the between-platform MBH and multi-omics QTL analyses that a narrower study could not attempt.
- The cohort's multiethnic composition (predominantly Arab, Filipino, and Indian ancestry) is a genuine strength relative to the largely Caucasian-focused prior multiomics literature the authors cite, even though it does not by itself establish broader generalizability (Section 7).
- Significance thresholds are explicitly scaled to the number of comparisons within each platform pair or comparison family, rather than a single threshold applied uniformly across very differently sized platforms.
- Non-genomics source data, R-loading scripts, and the COmics server source code (with a Docker image) are made publicly available, which is what allowed me to trace the platform-pair sample overlaps and edge-type counts cited throughout this analysis.
- The paper reports negative or qualified results rather than omitting them — for example, the ~28% MBH non-detection rate between overlapping proteomics platforms, and the explicit acknowledgment that some TWAS findings may reflect T2D-related metabolic differences in the cohort rather than a novel mechanism.

8.2 Limitations

Acknowledged directly by the authors: population stratification risk from the multiethnic cohort (addressed with genotype principal components); non-fasting sample collection as a source of biological variability

(though a prior QMDiab study is cited as evidence this variability is random rather than systematically biasing); the possibility that some associations, including part of the TWAS lipid signature, reflect known T2D-associated metabolic shifts rather than new biology; and the acknowledgment that future work should extend to additional biofluids and matrices (secretome, exhalome, microbiome) not covered here.

From my own reading: the MBH pipeline is not described as covariate-adjusted in the way the GGM pipeline is (Section 4.2), and I could not determine from the material available whether this was addressed in supplementary methods. The moTWAS pipeline does adjust for several measured leukocyte populations (CD4 and CD8 T cells, NK cells, B cells, monocytes, granulocytes, and eosinophils), so it would not be accurate to call it unadjusted for immune-cell composition — but basophil abundance is not listed as a separate covariate, which matters specifically because four of the five transcripts dominating the reported TWAS-lipid signal are described as predominantly basophil-expressed; residual confounding from that one cell type specifically cannot be ruled out from the reported analysis (Section 6). The COMics interface does not visually distinguish edges by evidence type or effect size (Section 4.1), which is a usability and interpretability limitation for anyone using the tool without having first read the primary Methods section closely. The subtype-network provenance question in Section 7 could not be resolved from the material provided to me. None of these, individually, undermines the paper's central claim that broad multiomics integration is technically achievable and can recover known biology; together, they bound how far any single edge in the network — particularly a novel one without an independent replication or a genetic anchor — should be trusted without further checking.

Study feature	Scientific strength	Interpretive limitation	Useful next test
Cohort design (cross-sectional, non-fasting, case-control)	Cases and controls processed identically, minimizing batch effects between groups	Non-fasting design adds biological variability; cross-sectional design cannot establish temporal order for non-genetic layers	Compare a subset of key associations against a fasting replication sample, if available
Multiethnic composition (Arab, Filipino, Indian)	Broader ancestry representation than much prior multiomics work, which is largely Caucasian-focused	Still a single-country, single-site sample; does not establish generalizability beyond these three ancestry groups	Test whether key network edges replicate in an independent, differently composed cohort
Platform breadth (18 platforms)	Enables between-platform MBH and cross-omics QTL analyses not possible with fewer platforms	Different platforms measure very different numbers of traits, complicating direct comparison of edge density	Report per-platform statistical power alongside edge counts
MBH integration	Hypothesis-free, borrowed from a well-established bioinformatics concept (BBH/orthology)	Computed on raw, non-covariate-adjusted data unlike the paired GGM pipeline (Section 4.2)	Recompute a subset of MBH edges on covariate-residualized data and compare (Extension 1)
Association analyses (GWAS/EWAS/TWAS)	Genome-wide significance thresholds;	TWAS models do not separately adjust for	Explicitly test sensitivity of TWAS-lipid associations to

Study feature	Scientific strength	Interpretive limitation	Useful next test
	genetic anchor for GWAS supports reverse-causation control; TWAS models adjust for several major leukocyte populations	basophil abundance, the cell type most linked to the top TWAS genes	basophil abundance (Extension 2)
Diabetes-subtype networks	Extends prior subtype-signature work with a multiomic molecular context	Subgroup provenance within QMDiab could not be verified from material available (Section 7)	Clarify subgroup assignment method and per-subgroup sample sizes in supplementary material
COMics resource	Open, no-code access to a large integrated network for hypothesis generation	Edges are not visually distinguished by evidence type or strength in the interface (Section 4.1)	Add an edge-type/evidence-strength legend or filter to the COMics interface

Table 2. Balanced evaluation of key study features: strengths, interpretive limitations, and a useful next test for each.

9. Reproducibility

Non-genomics source data for every platform are deposited on Figshare in both Excel and R (.rda) formats, with an accompanying R script for loading the .rda files, and the source code and a Docker image for the COMics server are available on GitHub with an archived Zenodo DOI. This is a genuinely reproducible-in-principle setup for the non-genomics side of the analysis. Genetic and methylation data are explicitly not deposited publicly because the informed consent obtained from participants does not cover public posting of individual-level genotype or methylation data; researchers affiliated with a research institution may request access from the corresponding authors subject to institutional review board approval. The restricted genotype and methylation data prevent complete independent reproduction of analyses requiring those layers (the eQTL, meQTL, GWAS, and EWAS pipelines) without approved access. The moTWAS (RNA-expression-omics) pipeline is a partial exception worth separating out: transcriptomics data are among the platforms described as part of the non-genomics Figshare deposit, so the extent to which this specific analysis can be reproduced depends on whether the released expression files, the participant-matched omics files, and the leukocyte-proportion covariates (CD8, CD4, NK, Bcell, Mono, Gran, Eos) used in the model are all present with matching participant identifiers — I could not confirm this from the materials available to me, so I would not say categorically that moTWAS cannot be reproduced, only that I cannot currently confirm that it can be.

The Methods section names specific software and package versions for several stages (R version 4.1.0 and above, RStudio 1.4.1717 and above, the GeneNet R package for partial correlations, PLINK version 1.9). From the materials available to me, I could not determine whether the exact k-nearest-neighbors imputation parameters used before GGM computation, or the specific random seed (if any) affecting the KNN imputation, are documented precisely enough for an external researcher to reproduce the identical 6,183 GGM edges rather than a closely similar set. Without rerunning the analysis, I cannot determine whether different imputation settings would affect only a small number of edges or produce more substantial changes in the GGM network.

10. My Critical Perspective

I find the paper's central technical claim convincing: 18 platforms applied across the same 391-participant cohort (with coverage varying by platform, and platform-pair overlap of 229–356 participants) can be integrated into one network, and at least some of the recovered relationships — the caffeine pathway, the thyroxine-SERPINA7 binding relationship, the APOE-lipid associations, the ST3GAL1-glycan GWAS hit — line up well with independently established biology. That is a real and useful demonstration, and the fact that the authors made the underlying data and code available is what let me trace these claims back to specific numbers rather than take them on faith.

What I am less convinced of is that the network, as visualized and released through COMics, should be read by an end user as a single coherent evidence base. Before working through the Methods section, I would have treated all edges in the network as broadly comparable units of “association”; the methods show why that would be misleading, both because GWAS, EWAS, TWAS, GGM, and MBH edges rest on different causal footing (Sections 2.3 and 6), and because even within the network, the GGM and MBH pipelines were not preprocessed identically with respect to covariate adjustment (Section 4.2). I also think the diabetes-subtype networks, while a reasonable extension of the platform, rest on cross-cohort signature definitions that I could not fully trace through the material available to me, and should be read as specific to this cohort and pipeline rather than as generalizable T2D-subtype biology.

The authors themselves describe their associations as “hypothesis-generating in nature,” and I think that framing is the right one to hold onto throughout — not because the work is weak, but because a network this large, built from several statistically and biologically different kinds of evidence at different levels of causal strength, is exactly the kind of resource where the value is in generating specific, testable leads rather than in treating any single edge as a settled finding.

11. What Would Change My Mind

- If the COMics interface (or a future version of it) visually distinguished edges by evidence type and, where available, effect size or replication status, rather than only by the omics category of the connected nodes.
- If a sensitivity analysis showed that MBH edges computed on covariate-residualized data reproduce substantially the same network as the MBH edges computed on raw data reported here — that would resolve the concern in Section 4.2 directly.
- If the supplementary material clarified exactly how the MARD/MOD/SIRD/SIDD subgroup assignments were derived within QMDiab, and how many participants fall into each subgroup.
- If a handful of the paper's novel (not literature-replicated) MBH or TWAS findings were independently tested in a second cohort with basophil-proportion data specifically available, to check whether the TWAS-lipid signal in Section 6 survives adjustment for that one cell type beyond the leukocyte populations the published models already control for.

12. Potential Research Extensions

Extension 1 — Covariate-adjusted MBH as an alternative edge-selection rule

Limitation: MBH edges are computed on raw, unadjusted data while GGM edges are computed on covariate-residualized data (Section 4.2). Research question: how much does the MBH edge set change if the same age/sex/BMI/diabetes-state residualization used for GGM is applied before computing reciprocal Spearman correlations? Method: using the QMDiab non-genomics data available on Figshare, reconstruct a small, clearly bounded part of the network (for example, the metabolomics-to-lipidomics MBH edges) under both the published raw-data pipeline and a residualized version, and compare the resulting edge lists. Evidence generated: an overlap count and a list of edges present under one preprocessing choice but not the other. Interpretation: a largely overlapping edge set would suggest the confounding concern in Section 4.2 has limited practical impact; a substantially different edge set would indicate the network's between-platform edges are more sensitive to preprocessing choices than the single-pipeline result suggests.

Extension 2 — Basophil-specific sensitivity of the TWAS lipid signal

Limitation: the published moTWAS models already adjust for CD4 and CD8 T cells, NK cells, B cells, monocytes, granulocytes, and eosinophils, but not separately for basophils — and four of the five transcripts dominating the reported lipid associations (GATA2, HDC, MS4A3, FCER1A) are described as predominantly basophil-expressed (Section 6). Research question: does the association between these four transcripts and their associated lipid/lipoprotein traits persist after explicitly adding an estimated basophil proportion as a covariate, on top of the leukocyte populations already in the published model? Method: if public reference expression signatures for blood cell types are available, apply a standard cell-type deconvolution method to estimate basophil proportion from bulk expression, then re-test the transcript-lipid associations with that estimate added to the existing covariate set, using the public QMDiab expression and lipidomics data. Evidence generated: association statistics with and without the added basophil-proportion term, for the affected trait pairs. Interpretation: associations that remain significant after adding basophil proportion would strengthen the direct-transcription interpretation; associations that are substantially attenuated would support residual basophil confounding specifically, over and above the leukocyte populations the authors already controlled for.

Extension 3 — Reconstructing a small part of the COMics network and inspecting edge provenance

This is the project I would most like to attempt as a next step in my own training, using the authors' publicly available data and code rather than requesting restricted genotype or methylation access. I would take a single, clearly bounded molecule (for example, one metabolite with a modest number of reported MBH partners, so the reconstruction stays tractable) and rebuild its local neighborhood in the network from the public Figshare data, checking each edge's evidence type (GGM, MBH, or a reported genomics association) against what is stated in the Methods. I would then compare the published MBH selection for that molecule against one alternative edge-selection rule — for instance, the top correlated partner without the reciprocity requirement, or the reciprocal best hit computed after the same covariate residualization used for GGM — and note where the two approaches agree and where they diverge. I am presenting this as a project I would like to attempt, not as work I have already completed; I do not yet know how much the two approaches would differ in practice, and that uncertainty is exactly why I think the comparison would be informative.

13. Questions the Paper Leaves Open

- How much would the MBH edge set change if the reciprocal-correlation step were computed on covariate-residualized data instead of raw data, matching the GGM pipeline?
- The published moTWAS models adjust for several leukocyte populations but not for basophils specifically — how much of the residual TWAS association between GATA2/HDC/MS4A3/FCER1A and circulating lipids, after those existing adjustments, might still reflect basophil proportion differences between participants?
- Were the MARD/MOD/SIRD/SIDD subgroup assignments for QMDiab participants derived within this study, or imported directly from the earlier Qatar Biobank subtype-definition study, and how many QMDiab participants fall into each subgroup?
- Of the roughly 2,100 MBH edges reported network-wide, what fraction connect platforms measuring genuinely overlapping molecule sets (where an identity-like interpretation is plausible) versus platforms measuring distinct molecule classes (where a coupling or pathway interpretation is more appropriate)? The paper reports concordance rates only for three specific overlapping-platform pairs.

14. Lessons Learned

A few things changed how I read multiomics network papers after working through this one. First, a single network visualization can flatten several genuinely different levels of statistical evidence into one visual grammar — before this analysis I would not have thought to check, edge by edge, what correction scheme and covariate set each part of a combined network was held to; I now think that is one of the first things worth checking in any integrative multiomics resource, not an afterthought.

Second, illustrative validation cases like the caffeine and HbA1c vignettes are reassuring about a method's sensitivity to known biology, but I had been treating that kind of demonstration as closer to a general validation of the network's overall precision than it actually is. Those are two different questions, and a paper can answer the first convincingly without saying much about the second, and without necessarily having designed the examples as formal positive controls in the first place.

Third, I had not previously thought carefully about how whole-blood transcriptomic associations can be confounded by specific immune-cell types, as opposed to immune-cell composition in general. My first pass through this paper treated the moTWAS pipeline as unadjusted for immune-cell composition; going back to the exact model formula in the Methods section showed that it already adjusts for CD4/CD8 T cells, NK cells, B cells, monocytes, granulocytes, and eosinophils. That was a useful reminder to check the specific covariate list rather than assume a category is either “adjusted” or “unadjusted” from a general impression. What remains is a narrower point: basophils specifically are not in that covariate list, and four of the five key TWAS genes are predominantly basophil-expressed. That is now something I want to check for specifically — not just whether a model adjusts for cell composition, but which cell types it does and does not include — whenever I see “whole blood” as the tissue source for an expression-based association.

15. Connection to My Development as a Bioinformatician

Working through this paper's Methods section in enough detail to write Sections 4 and 6 above was, honestly, more work than reading the Results and Discussion — and that gap is itself the main thing this paper motivates

me to develop further. Specific skills I want to build, in roughly the order I would need them for the project in Extension 3:

- Multiomics data preprocessing — normalizing and residualizing heterogeneous data types (metabolomics, proteomics, methylation) on comparable scales before any correlation or association step.
- Statistical network construction — implementing partial-correlation (GGM) and reciprocal best-hit (MBH) pipelines myself, rather than only reading about them, so I understand their assumptions and failure modes directly.
- Dimensionality reduction and integration — understanding how principal components (genotype PCs, protein PCs) are used as covariates to absorb structure like population stratification before testing for associations.
- Multiple-testing correction — applying comparison-family-specific Bonferroni thresholds correctly when combining results across platforms of very different sizes.
- Network stability analysis — testing how sensitive a reconstructed network is to reasonable variations in preprocessing, missing-data handling, or edge-selection rules, rather than treating a single published pipeline's output as fixed.
- Reproducible workflows in R — since the QMDiab data and loading scripts are provided in R, building fluency there rather than translating everything into Python first.
- Biological annotation of computational results — connecting a statistical edge (like GATA2-lipid) to what is actually known about that gene's expression pattern (basophil-specific) before interpreting the edge biologically.
- Critical interpretation of high-dimensional associations — keeping the distinction between statistical significance, biological plausibility, and demonstrated mechanism explicit rather than letting them blur together, especially in a network this large.

The concrete next step I would like to attempt is the project described in Extension 3: using the authors' publicly available data and code to reconstruct a small, clearly defined part of the COMICS network, inspect the provenance of each edge in that neighborhood, and compare the published MBH edge-selection rule against one alternative rule. I do not have the applied experience yet to know exactly how that comparison would come out, but working through this paper's methods in detail has given me a specific enough understanding of what MBH and GGM are actually doing that I think I know what a meaningful comparison would look like, which feels like the right place to start from.

16. Conclusion

Halama et al. convincingly demonstrate that 18 molecular phenotyping platforms can be integrated across a deeply characterized cohort of 391 participants, although participant coverage differs across platforms and individual platform pairs share between 229 and 356 samples, and that at least a subset of the recovered relationships — the caffeine pathway, the thyroxine-SERPINA7 link, the ST3GAL1-glycan GWAS association — line up with independently established biology. The paper is also commendably transparent about several of its own limitations: the ~28% MBH miss rate between overlapping proteomics platforms, the basophil-specific expression of the genes driving the largest TWAS signal, and the possibility that some findings reflect known T2D metabolic shifts rather than new biology.

It does not establish that every edge in the network carries the same evidentiary weight, since GGM, MBH, GWAS, EWAS, and TWAS edges rest on different causal footing and, in the case of GGM versus MBH specifically, were not preprocessed under the same covariate-adjustment standard. It also does not establish, from the material available to me, the internal provenance of the diabetes-subtype networks, or that any individual novel association (as opposed to the illustrative validation cases) has been independently validated beyond the statistical association reported here.

The strongest evidence in the paper, from my reading, is the combination of the caffeine/HbA1c biological sanity checks and the genetically anchored ST3GAL1 finding, both of which provide a form of validation that does not depend solely on the internal consistency of the network itself. The most useful next step, both for evaluating this paper further and for my own development, would be a direct comparison of the MBH edge set under the paper's raw-data pipeline against a covariate-adjusted alternative — a comparison that is feasible with the publicly available data and that I would like to attempt as described in Section 12.

References

Halama, A., Zaghlool, S., Thareja, G., Kader, S., Al Muftah, W., Mook-Kanamori, M., Sarwath, H., et al. (2024). A roadmap to the molecular human linking multiomics with population traits and diabetes subtypes. *Nature Communications*, 15, 7111. <https://doi.org/10.1038/s41467-024-51134-x>

Ahlqvist, E., Storm, P., Karäjamäki, A., Martinell, M., Dorkhan, M., Carlsson, A., Vikman, P., et al. (2018). Novel subgroups of adult-onset diabetes and their association with outcomes: a data-driven cluster analysis of six variables. *The Lancet Diabetes & Endocrinology*, 6(5), 361–369. (Cited within the reviewed paper as the source of the MARD/MOD/SIRD/SIDD subtype framework discussed in Section 7.)

Zaghlool, S. B., Sharma, S., Molnar, M., Matias-Garcia, P. R., Elhadad, M. A., Waldenberger, M., Peters, A., et al. (2022). Metabolic and proteomic signatures of type 2 diabetes subtypes in an Arab population. *Nature Communications*, 13, 7121. (Cited within the reviewed paper as the source of the diabetes-subtype protein and metabolite signatures used to build the subgroup networks discussed in Section 7.)

Krumsiek, J., Suhre, K., Illig, T., Adamski, J., & Theis, F. J. (2011). Gaussian graphical modeling reconstructs pathway reactions from high-throughput metabolomics data. *BMC Systems Biology*, 5, 21. (Cited within the reviewed paper as the methodological basis for the GGM approach discussed in Sections 2.1 and 4.2.)

Overbeek, R., Fonstein, M., D'Souza, M., Push, G. D., & Maltsev, N. (1999). The use of gene clusters to infer functional coupling. *Proceedings of the National Academy of Sciences*, 96, 2896–2901. (Cited within the reviewed paper as one basis for the bidirectional-best-hit / mutual-best-hit concept discussed in Section 2.2.)

Karlsson, M., Zhang, C., Mear, L., Zhong, W., Digre, A., Katona, B., Sjöstedt, E., et al. (2021). A single-cell type transcriptomics map of human tissues. *Science Advances*, 7, eabh2169. (Cited within the reviewed paper as the source for the basophil-specific expression pattern discussed in Section 6.)